


Lecture 1.0

Reinforcement Learning: MDPs

Alexey Gruzdev
alexey.s.gruzdev@gmail.com
HSE, Winter 2020

RL Mechanics



Major Components of an RL Agent

- An RL agent may include one or more of these components:
 - Policy: agent's behavior function
 - Value function: how good is each state and/or action
 - Model: agent's representation of the environment

Policy

- A **policy** is the agent's behavior
- It is a map from state to action, e.g.
- Deterministic policy: $a = \pi(s)$
- Stochastic policy: $\pi(a \mid s) = P[A_t = a \mid S_t = s]$

Value Function

- Value function is a prediction of future reward
- Used to evaluate the goodness/badness of states
- And therefore to select between actions, e.g.

$$v_{\pi}(s) = E_{\pi}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t = s]$$

Model

- A **model** predicts what the environment will do next
- P predicts the next state
- R predicts the next (immediate) reward, e.g.

$$P_{ss'}^a = P[S_{t+1} = s' \mid S_t = s, A_t = a]$$
$$R_s^a = E[R_{t+1} \mid S_t = s, A_t = a]$$

Rewards hypothesis revisited

- A **reward** R_t is a scalar feedback signal
- Indicates how well agent is doing at step t
- The agent's job is to maximize cumulative reward

Reinforcement Learning is based on the **reward hypothesis**.

The reward hypothesis:

All goals can be described by the maximization of expected cumulative reward .

- Online banners recommender system
- Personal mobile-phone assistants

Markov Processes Family

- Markov Processes (Markov Chain)
- Markov Reward Processes
- Markov Decision Processes
- Extensions to MDPs:
 - Infinite & Continuous MDP
 - POMDP
 - Undiscounted MDP

Introduction to MDPs

- *Markov decision processes* formally describe an environment for reinforcement learning
- Where the environment is *fully observable*
- i.e. The current *state* completely characterizes the process
- Almost all RL problems can be formalized as MDPs, e.g.
 - Optimal control primarily deals with continuous MDPs
 - Partially observable problems can be converted into MDPs
 - Bandits are MDPs with one state

Markov Property

- Definition: a state S_t is **Markov** if and only if

$$P[S_{t+1} | S_t] = P[S_{t+1} | S_1, \dots, S_t]$$

- “The future is independent of the past given the present”
 - The state captures all relevant information from the history
 - Once the state is known, the history may be thrown away
 - i.e. The state is a sufficient statistic of the future

State Transition Matrix

- For a Markov state s and successor state s' , the *state transition probability* is defined by

$$P_{ss'} = P [S_{t+1} = s' \mid S_t = s]$$


- State transition matrix $P_{ss'}$ defines transition probabilities from all states s to all successor states s'

$$P_{ss'} = \begin{pmatrix} P_{11} & \dots & P_{1n} \\ \vdots & \ddots & \vdots \\ P_{n1} & \dots & P_{nn} \end{pmatrix}$$


Markov Process

- Markov process is a memoryless random process, i.e. a sequence of random states S_1, \dots, S_t with the Markov property.
- *A Markov Process (or Markov Chain)* is a tuple (S, P)
 - S is a (finite) set of states
 - $P_{ss'}$ is a state transition probability matrix
 - $P_{ss'} = P[S_{t+1} = s' \mid S_t = s]$

Markov Process: Student Example



Markov Process: Episodes Sampling




- Sample episodes for Student Markov Chain starting from

$$S_1 = C1$$

$$S_1, \dots, S_t$$

- C1 C2 C3 Pass Sleep
- C1 FB FB C1 C2 Sleep
- C1 C2 C3 Pub C2 C3 Pass Sleep
- C1 FB FB C1 C2 C3 Pub C1 FB
FB FB C1 C2 C3 Pub C2 Sleep

Markov Process: Transition Probabilities




$$\mathcal{P} = \begin{matrix} & \begin{matrix} C1 & C2 & C3 & Pass & Pub & FB & Sleep \end{matrix} \\ \begin{matrix} C1 \\ C2 \\ C3 \\ Pass \\ Pub \\ FB \\ Sleep \end{matrix} & \begin{bmatrix} & & & & & 0.5 & \\ & 0.5 & & & & & 0.2 \\ & & 0.8 & & & & \\ & & & 0.6 & 0.4 & & 1.0 \\ 0.2 & 0.4 & 0.4 & & & & \\ 0.1 & & & & & 0.9 & \\ & & & & & & 1 \end{bmatrix} \end{matrix}$$

Markov Reward Process

- A Markov reward process is a Markov chain with values.
- A *Markov Reward Process* is a tuple (S, P, R, γ)
 - S is a (finite) set of states
 - $P_{ss'}$ is a state transition probability matrix
 - $P_{ss'} = P[S_{t+1} = s' \mid S_t = s]$
 - R is a reward function, $R_s = E[R_{t+1} \mid S_t = s]$
 - γ is a discount factor, $\gamma \in [0, 1]$

MRP: Student Example



Return

- The *return* G_t is the total discounted reward from time-step t .

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- The *discount* $\gamma \in [0, 1]$ is the present value of future rewards
- The value of receiving reward R after $k + 1$ time-steps is $\gamma^k R$.
- This values immediate reward above delayed reward.
 - γ close to 0 leads to "myopic" evaluation
 - γ close to 1 leads to "far-sighted" evaluation

Discount intuition

Most Markov reward and decision processes are discounted. Why?

- Mathematically convenient to discount rewards
- Avoids infinite returns in cyclic Markov processes
- Uncertainty about the future may not be fully represented
- If the reward is financial, immediate rewards may earn more interest than delayed rewards
- Animal/human behavior shows preference for immediate reward
- It is sometimes possible to use *undiscounted* Markov reward processes (i.e. $\gamma = 1$), e.g. if all sequences terminate

Value Function

- The value function $v(s)$ gives the long-term value of state s
- The *state value function* $v(s)$ of an MRP is the expected return starting from state s

$$v(s) = E[G_t \mid S_t = s]$$

Student MRP returns

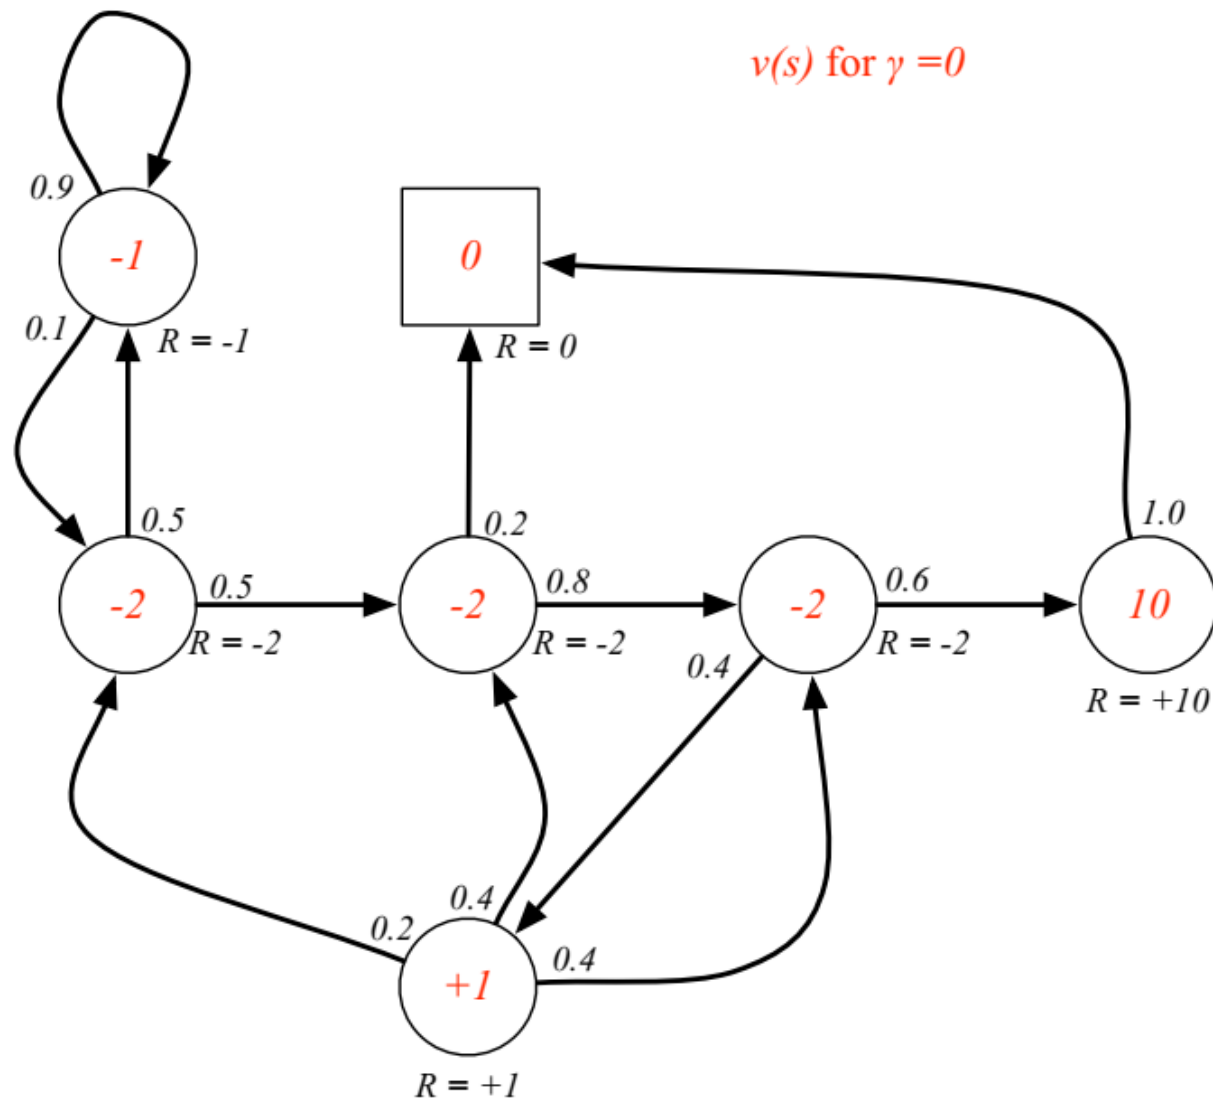
Sample returns for Student MRP with:

- $S_1 = C1$
- $\gamma = 0.5$


$$G_1 = R_2 + \gamma R_3 + \dots + \gamma^{T-2} R_T$$

C1 C2 C3 Pass Sleep	$v_1 = -2 - 2 * \frac{1}{2} - 2 * \frac{1}{4} + 10 * \frac{1}{8}$	=	-2.25
C1 FB FB C1 C2 Sleep	$v_1 = -2 - 1 * \frac{1}{2} - 1 * \frac{1}{4} - 2 * \frac{1}{8} - 2 * \frac{1}{16}$	=	-3.125
C1 C2 C3 Pub C2 C3 Pass Sleep	$v_1 = -2 - 2 * \frac{1}{2} - 2 * \frac{1}{4} + 1 * \frac{1}{8} - 2 * \frac{1}{16} \dots$	=	-3.41
C1 FB FB C1 C2 C3 Pub C1 ...	$v_1 = -2 - 1 * \frac{1}{2} - 1 * \frac{1}{4} - 2 * \frac{1}{8} - 2 * \frac{1}{16} \dots$	=	-3.20
FB FB FB C1 C2 C3 Pub C2 Sleep			

State-Value function for student MRP




State-Value function for student MRP



Markov Decision Process

- A Markov decision process (MDP) is a Markov reward process with decisions. It is an *environment* in which all states are Markov.
- A Markov reward process is a Markov chain with values.
- A *Markov Decision Process* is a tuple (S, A, P, R, γ)
 - S is a (finite) set of states
 - A is a finite set of actions
 - $P^a_{ss'}$ is a state transition probability matrix
 - $P^a_{ss'} = P[S_{t+1} = s' \mid S_t = s, A_t = a]$
 - R is a reward function, $R^a_s = E[R_{t+1} \mid S_t = s, A_t = a]$
 - γ is a discount factor, $\gamma \in [0, 1]$

Markov Decision Process: Example



MDP Policies

A *policy* π is a distribution over actions given states

$$\pi(a \mid s) = P[A_t = a \mid S_t = s]$$

- A policy fully defines the behavior of an agent
- MDP policies depend on the current state (not the history)
- i.e. Policies are *stationary* (time-independent),
 $A_t \sim \pi(\cdot \mid S_t); \forall t > 0$

MDP Policies

- Given an MDP $M = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$ and a policy π
- The state sequence S_1, \dots, S_t is a Markov process (S, P^π)
- The state and reward sequence S_1, R_1, S_2, \dots is a Markov reward process $(S, P^\pi, R^\pi, \gamma)$
- where

$$P^\pi_{s,s'} = \sum_{a \in \mathcal{A}} \pi(a | s) P^a_{s,s'}$$

$$R^\pi_s = \sum_{a \in \mathcal{A}} \pi(a | s) R^a_s$$

Updated Value functions


- The *state-value function* $v_{\pi}(s)$ of an MDP is the expected return starting from state s , and then following policy π

$$v_{\pi}(s) = E_{\pi} [G_t \mid S_t = s]$$

- The *action-value function* $q_{\pi}(s, a)$ is the expected return starting from state s , taking action a , and then following policy π

$$q_{\pi}(s, a) = E_{\pi} [G_t \mid S_t = s, A_t = a]$$

Student MDP: Value function



Optimal Value Function

- The *optimal state-value function* $v_*(s)$ is the maximum value function over all policies

$$v_*(s) = \max_{\pi} (v_{\pi}(s))$$

- The *optimal action-value function* $q_*(s; a)$ is the maximum action-value function over all policies

$$q_*(s, a) = \max_{\pi} (q_{\pi}(s, a))$$

- The optimal value function specifies the best possible performance in the MDP.
- An MDP is “solved” when we know the optimal value fn.

Optimal Policy

- Define a partial ordering over policies:

$$\pi \geq \pi' \text{ if } v_{\pi}(s) \geq v_{\pi'}(s) \quad \forall s$$

Theorem: *For any Markov Decision Process*

- *There exists an optimal policy π_* that is better than or equal to all other policies: $\pi_* \geq \pi \quad \forall \pi$*
- *All optimal policies achieve the optimal value function,*

$$v_{\pi_*}(s) = v_*(s)$$

- *All optimal policies achieve the optimal action-value function*

$$q_{\pi_*}(s, a) = q_*(s, a)$$


No more theory – let's do real things!




Cross Entropy method

- Let's derive some heuristic:
 - Initialize policy (random!)
 - Repeat:
 - Sample N episodes
 - Pick best 30 % best episodes – a.k.a. “elite” episodes
 - Change policy to choose actions from elite episodes


Cross Entropy method: step-by-step




Cross Entropy method: step-by-step




Cross Entropy method: step-by-step




Cross Entropy method: step-by-step




Cross Entropy method: step-by-step




Cross Entropy method: step-by-step




Cross Entropy method: step-by-step




Cross Entropy method: step-by-step




Cross Entropy method: step-by-step



Cross Entropy method: step-by-step



Cross Entropy method: step-by-step



Tabular Cross Entropy method

- Policy is matrix A
 - $\pi(a \mid s) = P[A_t = a \mid S_t = s] = A_{s,a}$
- Sample N sessions with that policy
- Get M best sessions (elites)
- Elite = $[(s_1, a_1), (s_2, a_2), \dots, (s_k, a_k)]$
- Update policy:

$$\pi(a \mid s) = \frac{\text{took } a \text{ at } s \text{ state}}{\text{was at } s \text{ state}} = \frac{\sum [s_t = s][a_t = a]}{\sum [s_t = s]}$$

Cross Entropy problems

But what if your environment has infinite/large state space ?



Approximated Cross Entropy method

- Policy is approximated
 - $\pi(a | s)$ predicted by Neural Network, Random Forest or any other ML algorithm.
- You can't set $\pi(a | s)$ explicitly, it's not matrix anymore.
- Training data for our model:

$$\text{Elite} = [(s_1, a_1), (s_2, a_2), \dots, (s_k, a_k)]$$

Questions?

